

Title - Landmark Detection

ABSTRACT

What is Landmark Detection?

The mechanism of detecting the famous human-made sculptures, buildings, and monuments inside an image is defined as Landmark Detection. You can simply compare it with the famous application of Google known as Google Landmark Detection, which is used by Google Maps.

At the end of this project, you will be able to create your own landmark detector like Google using the Keras library of Deep learning.

Dataset

We choose Kaggle's dataset for our deep learning project. The dataset consists of image URLs that are publicly available online. The dataset contains three CSV files including test images, training images, and index images. The test images are for the purpose of image recognition and landmark labelling predicted by the deep learning model. The training images are already defined and associated with landmark labels that are used to train models for accurate landmark recognition. The use of index images is found in the image retrieval task. You can download the dataset from

[URL: for dataset\(105gb\)](#)

<https://www.kaggle.com/competitions/landmark-recognition-2020/data>.

OBJECTIVE

Our task is to build neural networks to recognize the landmarks inside the images using the Python programming language. The most critical task for any project is to choose an appropriate dataset for model training.

This technology can predict landmark labels directly from image pixels, to help people better understand and organize their photo collections.

INTRODUCTION

Object recognition is one of the fundamental problems widely studied in computer vision. The problem that we will be investigating falls into this very category: we want to recognize landmarks (e.g. White House, Great Wall of China etc.) present in an image directly from the image pixels. Landmark recognition is interesting because it can help people better understand and organize their digital photo collections.

Despite being a straightforward objective, i.e. identifying landmark presented in a given image, the task itself is challenging especially as we increase the number of distinct landmarks. Our proposed landmark recognition problem can be best described as an instance-level recognition problem. It differs from categorical recognition problem in that instead of recognizing general entities such as mountains and buildings. Landmark recognition also differs from what we have seen in the ImageNet classification challenge. For example, since landmarks are generally rigid,

immobilized, one-of-the-kind objects, the intra-class variation is very small (in other words, a landmark's appearance does not change that much across different images of it). As a result, variations only arise due to image capture conditions, such as foreground/background clutter, occlusions, different viewpoints, weather and illumination, making this distinct from other image recognition datasets where images of a particular class (such as a cat) can vary much more in shape and appearance. These characteristics are also shared with other instance-level recognition problems, such as artwork recognition — ideally with mild modification, our solution to landmark recognition problem can be applied to research for other image recognition problems as well. In this project. landmark recognition problem.

METHODOLOGY

Data Collection

Data collection is the process of gathering and measuring information on variables of interest, in an established systematic fashion that enables one to answer stated research questions, test hypotheses, and evaluate outcomes.

Data collected from

<https://github.com/cvdfoundation/google-landmark>

The image displays two screenshots of the Kaggle website interface. The top screenshot shows the 'train.csv' dataset (37.07 MB) with a table of 1580470 unique values. The table has columns 'id', 'landmark_id', and 'count'. The bottom screenshot shows the 'Data Explorer' view for the same dataset, displaying 407 files as a grid of image thumbnails. Each thumbnail includes a unique ID and file size. The right sidebar in both screenshots shows the 'Data Explorer' panel with a file tree structure.

id	landmark_id	count
1580470 unique values		
17668ef415d37859	1	203k
92b629ed571448f6	1	
cd41bf948edc9348	1	
fb99f1e98cd2f78	1	
25c9dfc7ea69838d	7	
28b13f94a6f1f3c1	7	
387d6584f473ba35	7	
4a7ba9eb16d51bc4	7	
597353dfbb3df649	7	
a4898dc4fcc3a18	7	

Preprocessing

Data preprocessing, a component of data preparation, describes any type of processing performed on raw data to prepare it for another data processing procedure. It has traditionally been an important preliminary step for the data mining process.

- Resizing images
- Normalization- the process of organizing data in a database. This includes creating tables and establishing relationships between those tables according to rules designed both to protect the data and to make the database more flexible by eliminating redundancy and inconsistent dependency

augmentation

Data augmentation is a technique of artificially increasing the training set by creating modified copies of a dataset using existing data. It includes making minor changes to the dataset or using deep learning to generate new data points.

- Rotating
- Flipping
- Scaling
- Encoding

Model creation

- Import Python Libraries
- Read the Dataset
- Explore the Dataset
- Feature Selection
- Build the Model
- Evaluate the Model's Performance

Sequential model

Keras documentation the Sequential model is a linear stack of layers. You can create a Sequential model by passing a list of layer instances to the constructor: `from keras.models import Sequential from keras.layers import Dense, Activation model = Sequential([Dense(32, input_shape=(,)).we have the following declaration in it`

- Input layer
- output layer
- hidden layer
- CNN

A CNN can be instantiated as a Sequential model because each layer has exactly one input and output and is stacked together to form the entire network

- full connected layers

Fitting

Data fitting is the process of fitting models to data and analyzing the accuracy of the fit. Engineers and scientists use data fitting techniques, including mathematical equations and nonparametric methods, to model acquired data.

MATLAB lets you import and visualize your data, and perform basic fitting techniques such as polynomial and spline interpolation. You can perform data fitting interactively using the MATLAB Basic Fitting tool, or programmatically using MATLAB functions for fitting.

Split the data

Data splitting is when data is divided into two or more subsets. Typically, with a two-part split, one part is used to evaluate or test the data and the other to train the model. Data splitting is an important aspect of data science, particularly for creating models based on data. A commonly used ratio is 80:20, which means 80% of the data is for training and 20% for testing. We used 70:30 ie uses the first 70 percent of your data for training and the remaining 30 percent of the data for evaluation.

train model with data

The process of training an ML model involves providing an ML algorithm (that is, the *learning algorithm*) with training data to learn from. The term *ML model* refers to the model artifact that is created by the training process.

The training data must contain the correct answer, which is known as a *target* or *target attribute*. The learning algorithm finds patterns in the training data that map the input data attributes to the target (the answer that you want to predict), and it outputs an ML model that captures these patterns.

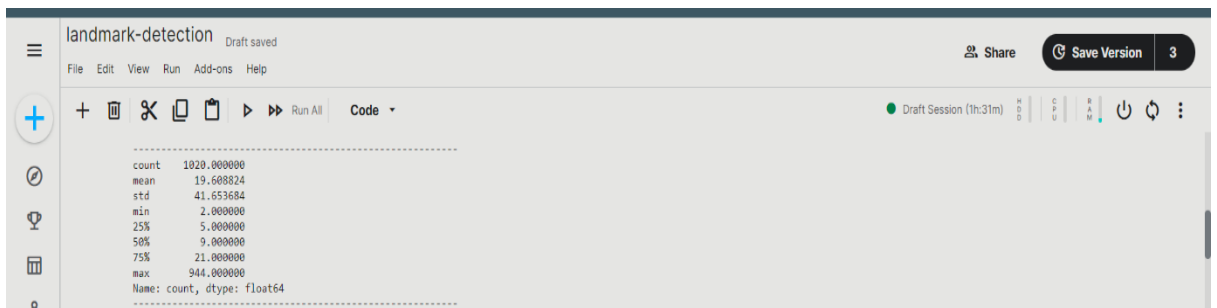
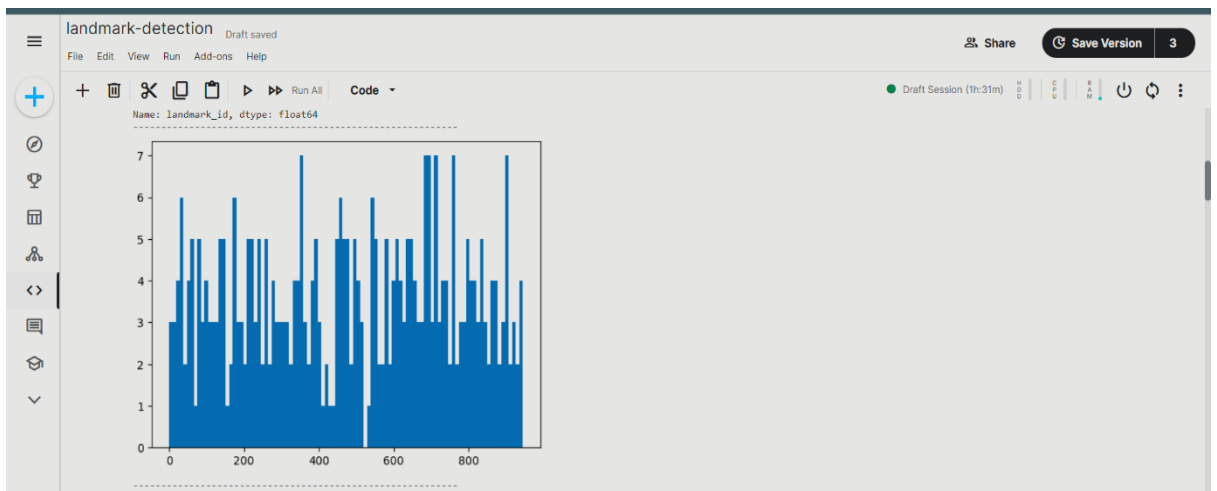
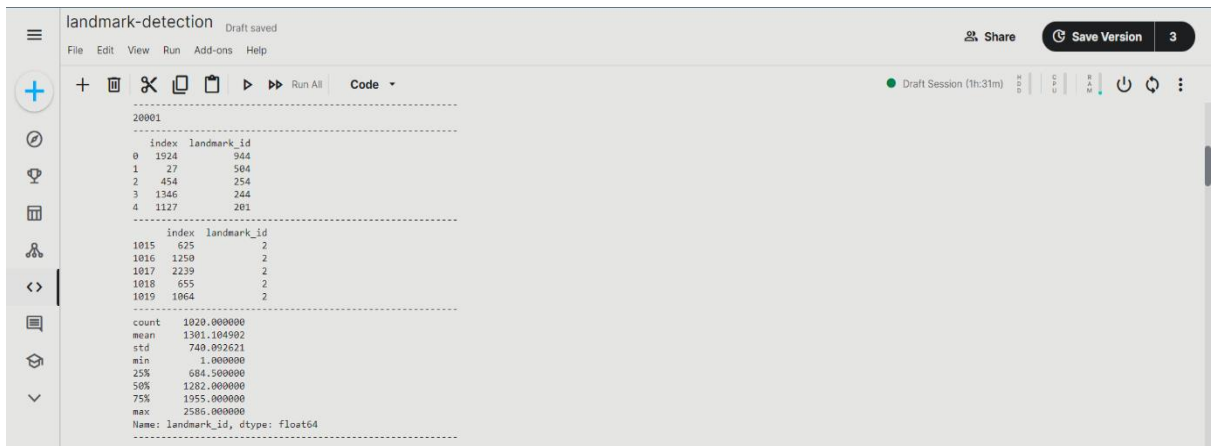
test the model-metric accuracy, 11 score


You build a model, get feedback from metrics, make improvements, and continue until you achieve a desirable classification accuracy. Evaluation metrics explain the performance of the model. An important aspect of evaluation metrics is their capability to discriminate among model results.

We perform prediction on the model. We can expect the output as show below

CODE

<https://www.kaggle.com/code/poornimamarini/landmark-detection>





landmark-detection Draft saved

File Edit View Run Add-ons Help

Share Save Version 3

Draft Session (1h:31m)

```
[50]:
print(df.head())
df=df.loc[df['id'].str.startswith('00',na=False).:]
print("-----")
num_classes=len(df["landmark_id"].unique())
print(num_classes)
print("-----")
num_data=len(df)
print(num_data)
print("-----")
data=pd.DataFrame(df["landmark_id"].value_counts())
data.reset_index(inplace=True)
print(data.head())
print(data.tail())
print("-----")
data.columns=["landmark_id","count"]
print(data["landmark_id"].describe())
print("-----")
plt.hist(data["landmark_id"],100,range=(0,61),label="test")
plt.show()
print("-----")
print(data["count"].describe())
print("-----")
plt.hist(data["count"],100,range=(0,61),label="test")
```

landmark-detection Draft saved

File Edit View Run Add-ons Help

Share Save Version 3

Draft Session (1h:31m)

```

plt.show()
print(data["count"].between(0,5).sum())
print("-----")
print(data["count"].between(5,10).sum())
print("-----")
plt.hist(df["landmark_id"],bins=df["landmark_id"].unique())
plt.show()

id landmark_id
119 00cbe0067c078490 27
120 00f928e383e1d121 27
796 009ecd56b5e9adb 60
1089 00d5d47528591d4 124
1113 00e0903e381ab009 134
-----
61
-----
79
-----
index landmark_id
0 1924 5
1 2202 3
2 454 3
3 392 2
4 1346 2
-----
index landmark_id
56 1236 1
```

landmark-detection Draft saved


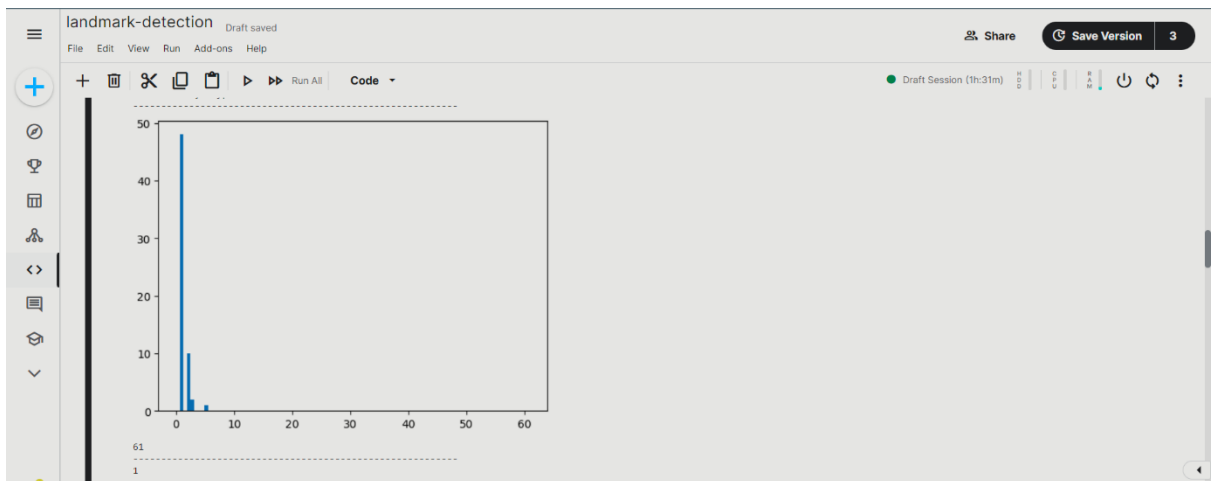
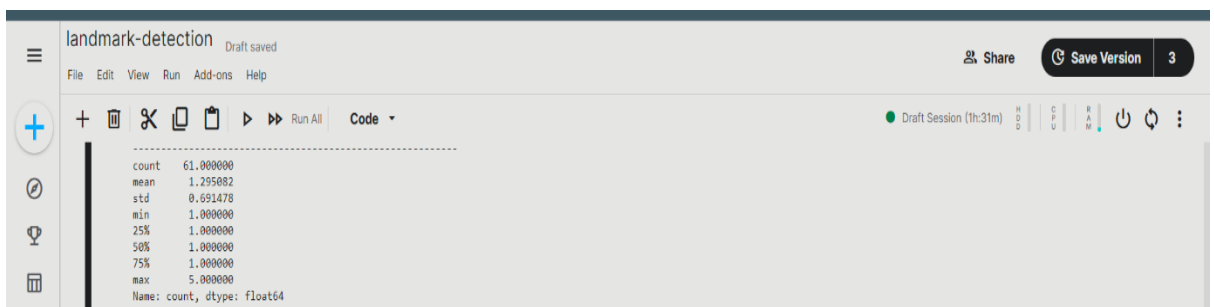

File Edit View Run Add-ons Help

Share Save Version 3

Draft Session (1h:31m)

```

index landmark_id
56 1236 1
57 1302 1
58 60 1
59 1393 1
60 2581 1
-----
count 61.000000
mean 1304.344202
std 733.481036
min 27.000000
25% 624.000000
50% 1359.000000
75% 1836.000000
max 2581.000000
Name: landmark_id, dtype: float64
-----
```



landmark-detection Draft saved

File Edit View Run Add-ons Help

Share Save Version 3

Draft Session (1h:32m)

+ Code + Markdown

```
[61]: print(lencoder.fit(df["landmark_id"]))
      print("-----")
      print(df.head())
      print("-----")

LabelEncoder()
-----
      id  landmark_id
119  00cba0067c078490    27
120  00f928e383e1d121    27
796  009ecd56b5e9adb     60
1089 00d5d47528839144   124
1133 00e9003a381ab809   134
-----

[12]: from sklearn.preprocessing import LabelEncoder
      lencoder = LabelEncoder()
      lencoder.fit(df["landmark_id"])
      def encode_label(lb1):
          return lencoder.transform(lb1)
      def decode_label(lb1):
          return lencoder.inverse_transform(lb1)
```

landmark-detection Draft saved

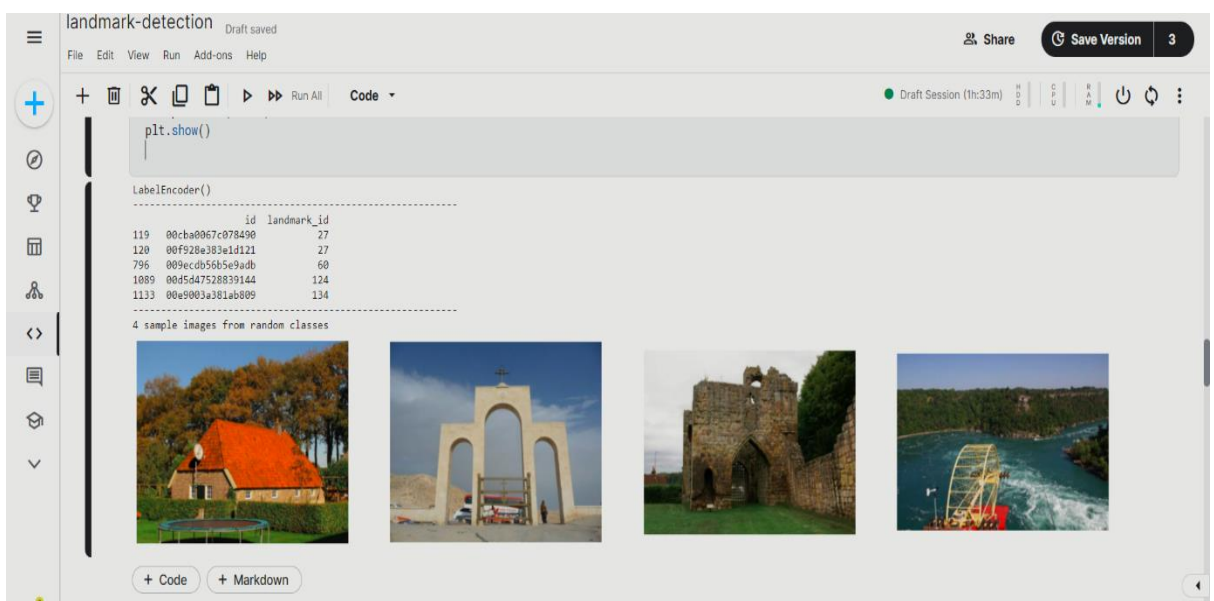
File Edit View Run Add-ons Help

Share Save Version 3

Draft Session (1h:32m)

+ Code + Markdown

```
[14]: from sklearn.preprocessing import LabelEncoder
      lencoder = LabelEncoder()
      lencoder.fit(df["landmark_id"])
      def encode_label(lb1):
          return lencoder.transform(lb1)
      def decode_label(lb1):
          return lencoder.inverse_transform(lb1)
      def get_image_from_number(num, df):
          fname, label = df.iloc[num, :]
          f1 = fname[0]
          f2 = fname[1]
          f3 = fname[2]
          path = os.path.join(f1, f2, f3, fname)
          im = cv2.imread(os.path.join(base_path, path))
          return im, label
      print("4 sample images from random classes")
      fig = plt.figure(figsize=(16, 16))
      for i in range(1, 5):
          ri = random.choices(os.listdir('/kaggle/input/landmark-recognition-2020/train/0/'), k=3)
          folder = base_path + '/' + ri[0] + '/' + ri[1] + '/' + ri[2]
          random_img = random.choice(os.listdir(folder))
          img = np.array(Image.open(folder + '/' + random_img))
          fig.add_subplot(1, 4, i)
          plt.imshow(img)
      plt.axis('off')
```



landmark-detection Draft saved

File Edit View Run Add-ons Help

Share Save Version 3

Draft Session (1h:33m)

build the model*

[62]:

```
from keras.applications import VGG19
from keras.layers import *
from keras import Sequential
# Parameters
# learning_rate = 0.0001
# decay_speed = 1e-6
# momentum = 0.09
# loss_function = "sparse_categorical_crossentropy"
source_model = VGG19(weights=None)
#new_layer = Dense(num_classes, activation=activations.softmax, name='prediction')
drop_layer = Dropout(0.5)
drop_layer2 = Dropout(0.5)
model = Sequential()
for layer in source_model.layers[:-1]: # go through until last layer
    if layer == source_model.layers[-25]:
        model.add(BatchNormalization())
        model.add(layer)
model.add(Dense(num_classes, activation = "softmax"))
print(model.summary())
```

landmark-detection Draft saved

File Edit View Run Add-ons Help

Share Save Version 3

Draft Session (1h:33m)

Model: "sequential_4"

Layer (type)	Output Shape	Param #
batch_normalization_2 (Batch Normalization)	(None, 224, 224, 3)	12
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv4 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160

landmark-detection Draft saved

File Edit View Run Add-ons Help

Share Save Version 3

Draft Session (1h:34m)

block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359008
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359008
block4_conv4 (Conv2D)	(None, 28, 28, 512)	2359008
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359008
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359008
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359008
block5_conv4 (Conv2D)	(None, 14, 14, 512)	2359008
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 4096)	102764544
fc2 (Dense)	(None, 4096)	16781312
dense_2 (Dense)	(None, 61)	249917

=====

Total params: 139,820,169

Trainable params: 139,820,163

Non-trainable params: 6

landmark-detection

Draft saved

Share

Save Version

3

File Edit View Run Add-ons Help

+ - ✂ 📄 📁 ▶ ⏮ Run All Code ▾

Draft Session (1h:34m)

CODE

RAM

⏻ ↺ ⋮

FITTING

[63]:

```
#Function used to process the data, fitted into a data generator.
def get_image_from_number(num, df):
    fname, label = df.iloc[num,:]
    fname = fname + ".jpg"
    f1 = fname[0]
    f2 = fname[1]
    f3 = fname[2]
    path = os.path.join(f1, f2, f3, fname)
    im = cv2.imread(os.path.join(base_path, path))
    return im, label
def image_reshape(im, target_size):
    print(im.shape)
    return cv2.resize(im, target_size)
def get_batch(dataframe, start, batch_size):
    image_array = []
    label_array = []
    end_img = start+batch_size
    if end_img > len(dataframe):
        end_img = len(dataframe)
    for idx in range(start, end_img):
```

landmark-detection

Draft saved

Share

Save Version

3

File Edit View Run Add-ons Help

+ - ✂ 📄 📁 ▶ ⏮ Run All Code ▾

Draft Session (1h:35m)

CODE

RAM

⏻ ↺ ⋮

```
r2 = tname[1]
f3 = fname[2]
path = os.path.join(f1, f2, f3, fname)
im = cv2.imread(os.path.join(base_path, path))
return im, label
def image_reshape(im, target_size):
    print(im.shape)
    return cv2.resize(im, target_size)
def get_batch(dataframe, start, batch_size):
    image_array = []
    label_array = []
    end_img = start+batch_size
    if end_img > len(dataframe):
        end_img = len(dataframe)
    for idx in range(start, end_img):
        n = idx
        im, label = get_image_from_number(n, dataframe)
        im = image_reshape(im, (224, 224)) / 255.0
        image_array.append(im)
        label_array.append(label)
    label_array = encode_label(label_array)
    return np.array(image_array), np.array(label_array)
```

landmark-detection

Draft saved

Share

Save Version

3

File Edit View Run Add-ons Help

+ - ✂ 📄 📁 ▶ ⏮ Run All Code ▾

Draft Session (1h:35m)

CODE

RAM

⏻ ↺ ⋮

split the data

[64]:

```
batch_size = 64
epoch_shuffle = True
weight_classes = True
epochs = 1
# Split train data up into 70% and 30% validation
train, validate = np.split(df.sample(frac=1), [int(0.7*len(df))])
print("Training on:", len(train), "samples")
print("Validation on:", len(validate), "samples")
```

Training on: 55 samples
Validation on: 24 samples

landmark-detection

Draft saved

Share

Save Version 3

File Edit View Run Add-ons Help

+

✂

📄

📄

▶

▶▶

Run All

Code

Draft Session (1h:35m)

⏮

⏪

⏩

⏭

⏹

[56]:

```

for e in range(epochs):
    print("Epoch: ", str(e+1) + "/" + str(epochs))
    if epoch_shuffle:
        train = train.sample(frac = 1)
        for it in range(int(np.ceil(len(train)/batch_size))):
            X_train, y_train = get_batch(train, it*batch_size, 100)

            model.train_on_batch(X_train, y_train)
        model.save("Model")

```

```

Epoch: 1/1
(800, 800, 3)
(600, 800, 3)
(482, 640, 3)
(800, 533, 3)
(600, 800, 3)
(281, 799, 3)
(600, 800, 3)
(533, 800, 3)
(533, 800, 3)
(600, 800, 3)
(800, 600, 3)
(548, 800, 3)
(536, 800, 3)
(800, 600, 3)
(600, 800, 3)
(600, 800, 3)

```

landmark-detection

Draft saved

Share

Save Version 3

File Edit View Run Add-ons Help

+

✂

📄

📄

▶

▶▶

Run All

Code

Draft Session (1h:35m)

⏮

⏪

⏩

⏭

⏹

TEST

[53]:

```

### Test on the training set
batch_size = 16
errors = 0
good_preds = []
bad_preds = []
for it in range(int(np.ceil(len(validate)/batch_size))):
    X_train, y_train = get_batch(validate, it*batch_size, batch_size)
    result = model.predict(X_train)
    cla = np.argmax(result, axis=1)
    for idx, res in enumerate(result):
        print("Class:", cla[idx], "- Confidence:", np.round(res[cla[idx]],2), "- GT:", y_train[idx])
        if cla[idx] != y_train[idx]:
            errors = errors + 1
            bad_preds.append([batch_size*it + idx, cla[idx], res[cla[idx]]])
        else:
            good_preds.append([batch_size*it + idx, cla[idx], res[cla[idx]]])
print("Errors: ", errors, "Acc:", np.round(100*(len(validate)-errors)/len(validate),2))
#Good predictions
good_preds = np.array(good_preds)
good_preds = np.array(sorted(good_preds, key = lambda x: x[2], reverse=True))

```

landmark-detection

Draft saved

Share

Save Version 3

File Edit View Run Add-ons Help

+

✂

📄

📄

▶

▶▶

Run All

Code

Draft Session (1h:35m)

⏮

⏪

⏩

⏭

⏹

```

cla = np.argmax(result, axis=1)
for idx, res in enumerate(result):
    print("Class:", cla[idx], "- Confidence:", np.round(res[cla[idx]],2), "- GT:", y_train[idx])
    if cla[idx] != y_train[idx]:
        errors = errors + 1
        bad_preds.append([batch_size*it + idx, cla[idx], res[cla[idx]]])
    else:
        good_preds.append([batch_size*it + idx, cla[idx], res[cla[idx]]])
print("Errors: ", errors, "Acc:", np.round(100*(len(validate)-errors)/len(validate),2))
#Good predictions
good_preds = np.array(good_preds)
good_preds = np.array(sorted(good_preds, key = lambda x: x[2], reverse=True))
fig=plt.figure(figsize=(16, 16))
for i in range(1,6):
    n = int(good_preds[i,0])
    img, lbl = get_image_from_number(n, validate)
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    fig.add_subplot(1, 5, i)
    plt.imshow(img)
    lbl2 = np.array(int(good_preds[i,1])).reshape(1,1)
    sample_cnt = list(df.landmark_id).count(lbl)
    plt.title("Label: " + str(lbl) + "nClassified as: " + str(decode_label(lbl2)) + "nSamples in class " + str(lbl) + ": " + str(sample_cnt))
    plt.axis('off')
plt.show()

```

- **CONCLUSION**

You can see the model output, how images are classified depending on the classes and labels. It uses the Keras library of deep learning to create a convolutional network which in turn trains the model. We perform good predictions results the above as Expected output for prediction. As the data has large memory(105gb) run it in Kaggle. we tackled the problem of Landmark recognition using transfer learning and data augmentation.